

Review Session

Prof. Alexandra Chouldechova
95-791: Data Mining

December 8, 2017

Agenda

- **Central themes**
- **Review of Supervised learning**
 - LDA, QDA
 - Regression, Classification methods list
 - Model selection approaches
 - Trees, Bagging, Random forests, Boosting
- **Review of Unsupervised learning**
 - Methods list
 - K -means, GMM's
 - Hierarchical clustering
 - Dimensionality Reduction
- **Overview: How to approach a particular task?**

Central themes

1. Generalizability

- We want to construct models that:
 - Capture **useful trends** in the data (don't *underfit*)
 - Ignore meaningless random fluctuations in the training data (don't *overfit*)
- To avoid **underfitting**:
 - Consider more flexible models
 - Try doing some feature engineering
- To avoid **overfitting**:
 - Use Cross-validation to identify models with low test error
 - Perform variable selection to reduce model variability
 - Try regularized models (with CV to select regularization parameter)

2. Bias-Variance tradeoff

- The *reducible* component of expected prediction error takes the form $\text{bias}(\hat{f})^2 + \text{var}(\hat{f})$
- Using a model with high **bias** risks *underfit*
- Using a model with with high **variance** risks *overfitting*

Central themes

3. Interpretability vs. Flexibility

- Highly structured models tend to be **easy** to interpret. Examples:
 - Linear and Logistic regression (with small p)
 - Regularized regression
 - Small decision trees
 - Additive models (with small p)
- Highly flexible models tend to do a better job at prediction, but are **hard** to interpret. Examples:
 - Large trees, Regression with large p
 - Bagging, Random Forests, Boosting

4. Feature engineering

- In practice, whether a model gives good results depends *a lot* on which features are fed into the model
- **Transforming** existing features can help performance *a lot*
 - Step functions, Polynomial regression, Splines, Additive models
- Data Visualization and Unsupervised learning techniques can help us to construct new features (e.g., PCA, Clustering)

Supervised learning methods: Regression

- Linear regression
- Polynomial regression, Step functions, Splines, Local regression
- Additive models
- Regularized regression (Lasso)
- k -Nearest Neighbours
- Decision trees
- Bagging
- Random forests
- Boosting

Supervised learning methods: Classification

- Logistic regression
 - Generalized Additive Models
 - Regularized logistic regression
- K -nearest neighbours
- Bayes methods:
 - Linear discriminant analysis
 - Quadratic discriminant analysis
 - Naive Bayes
- Classification trees
- Bagging
- Random forests
- Boosting

Model selection approaches

- **Cross-validation**
 - Calculate CV error for a bunch of models, and pick the one with the lowest CV error
 - Or: use the 1-SE rule
 - **Note:** If you are considering many different models, you should assess your final model on an unseen set of **test data**
- **Best subset selection, Forward stepwise selection** (to select variables)
 - Both methods give us the *best* k -predictor model for $k = 0, \dots, p$.
 - Can use AIC or BIC to select best model size
 - Best subset selection is computationally inefficient
 - Forward selection is OK, but Lasso generally works better if p isn't small
- **Regularized regression** (Lasso)
 - Automatically performs variable selection
 - Reduces model variance by shrinking estimated coefficients $\hat{\beta}_j$ towards 0

Bayes Methods

Bayes' theorem gives us a way of combining prior probabilities π_k , feature densities $f_k(x)$ and observed data to get a posterior probability that an observation with inputs x_0 is in class $y_0 = k$.

$$\hat{y}_0 = \operatorname{argmax}_{k=1,\dots,K} \frac{\pi_k f_k(x_0)}{\underbrace{\sum_{\ell=1}^K \pi_{\ell} f_{\ell}(x_0)}_{\text{posterior prob.}}} = \operatorname{argmax}_{k=1,\dots,K} \pi_k f_k(x_0)$$

- **Linear discriminant analysis (LDA)** assumes all the $f_k(x)$ are **Multivariate Normal** (μ_k, Σ)
 - Different means, same covariance matrix
- **Quadratic discriminant analysis (QDA)** assumes all the $f_k(x)$ are **Multivariate Normal** (μ_k, Σ_k)
 - Different means, different variances
- **Naive Bayes** assumes that within each class, all of the inputs are **independent**: $f_k(x) = \prod_{j=1}^p f_{k,j}(x_j)$

LDA vs. QDA

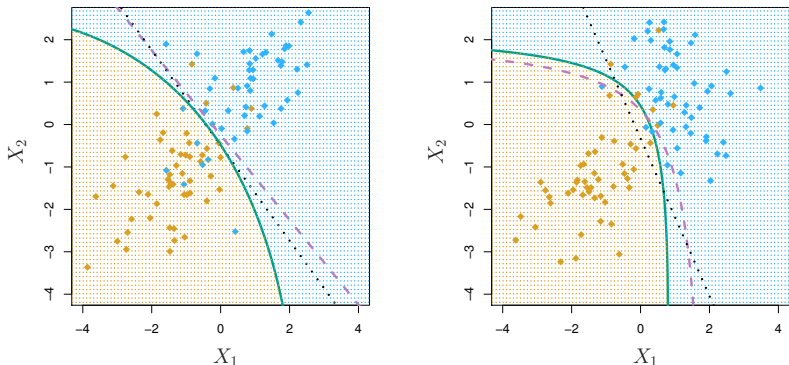




Figure 4.9 from ISL. Dashed purple curve is the Bayes classifier decision boundary. Solid green curve is QDA, dotted black line is LDA. Left: True boundary is linear. Right: True boundary is quadratic.

Naive Bayes vs. LDA vs. QDA


- Naive Bayes scales well to problems where p is large
 - If you have enough data to estimate the univariate density of each predictor (i.e., enough to form a *nice* histogram), you can apply Naive Bayes
- In LDA, we have to estimate $K \times p$ parameters to get $\hat{\mu}_k$'s and another $\frac{1}{2}p(p+1)$ parameters to estimate the $p \times p$ covariance matrix $\hat{\Sigma}$.
- In QDA, we have to estimate the means and K $p \times p$ covariance matrices. That's $\frac{1}{2}Kp(p+1)$ parameters!
- So why do even bother with methods like LDA or QDA?
 - They can capture meaningful **interactions**. Naive Bayes **cannot**.

Decision trees (CART)

- CART builds a deep tree via the **recursive binary partitioning** algorithm; the tree is then **pruned** to reduce variance
- At each step, one of the existing rectangles is split into two parts
- The splitting variable and splitting location is chosen to maximize the decrease in error (RSS for regression, Gini index for classification)
- Output: A tree representation, and a partition of the space into rectangles



Trees vs. Linear models



ISL Figure 8.7. Trees are bad when the boundary is linear, but very good when the boundary is well-described by a simple rectangular partition.

Classification trees vs. Other methods

Method	Interpretable	Flexible	Makes assumptions?
Logistic regression	Yes	Extensible	Yes
k -NN	No	Highly	No
LDA/QDA	Sometimes	No	Yes
Trees	Yes ¹	Somewhat	No

- Trees don't assume any particular relationship between the response Y and the inputs X_j , and large trees are quite flexible
- They tend to be poor predictors/classifiers because they are highly variable
 - Small changes in the training data can produce big changes in the resulting tree
- We can build ensembles of trees via Bagging, Random forests, and Boosting...but we lose interpretability

¹ Assuming the tree isn't too large

Bagging: Classification trees

- Given a training data (x_i, y_i) , $i = 1, \dots, n$, **bagging averages** the predictions from classification trees over a collection of bootstrap samples.
 - If we average over a bunch of high-variance, low-bias methods, we decrease the **variance**, while retaining **low bias**
- Here we'll describe how to apply Bagging to Classification Trees
 - ① For $b = 1, \dots, B$, get a bootstrap sample of size n from the training data: (x_i^{*b}, y_i^{*b}) , $i = 1, \dots, n$
 - ② Fit a classification tree $\hat{f}^{\text{tree}, b}$ on each sample
 - ③ Classify a new point x_0 by aggregating the votes (or probability estimates) across all B trees
- In Step (3), we can use the **Consensus** approach (classify to the class with the most votes), or the **Probability** approach (average the class probability estimates across all trees, classify to highest average probability)
- In Step (2) the trees are grown very large, with **no pruning**.

Random Forests

- **Random forests** provide an improvement over **bagged trees** by incorporating a small tweak that **decorrelates** the individual trees
 - This further **reduces variance** when we average the trees
- We still build each tree on a bootstrapped training sample
- But now, each time a split in a tree is considered, the tree may only split on a predictor from **a randomly selected subset of m predictors**
- A fresh selection of m randomly selected predictors is presented at each split... not for each tree, but for **each split of each tree**
- $m \approx \sqrt{p}$ turns out to be a good choice²
 - E.g., if we have 100 predictors, each split will be allowed to choose from among 10 randomly selected predictors

² $m = p/3$ is R's **randomForest** default for regression, $m = \sqrt{p}$ is the default for classification

Boosting algorithm: Regression trees

- ❶ Set $\hat{f}(x) = 0$ and *residuals* $r_i = y_i$ for all i in the training set
- ❷ For $b = 1, 2, \dots, B$, repeat:
 - ❶ Fit tree \hat{f}^b with d splits to the training data (X, r) ³
 - ❷ Update \hat{f} by adding *shrunk* version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$

- ❷ Update the residuals

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$$

- ❸ Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$

- λ is called the **shrinkage parameter**. Typically, $\lambda \ll 1$

³We're treating the residual vector r as our outcome at each step.

What's boosting trying to do?

- Boosting works best if d (the size of each tree) is small
- Given the current model, we fit a decision tree to the *residuals* from the model
- This new tree helps us perform just a little bit better in places where the current model wasn't doing well
- Unlike **bagging**, where each tree is large and tries to model the entire (bootstrapped) data well, each tree in **boosting** tries to incrementally improve the existing model
- Think of **boosting** as **learning slowly**: We use small trees, try to make incremental improvements, and further slow down the learning process by incorporating the *shrinkage parameter* λ

Random forests

- Uses bagging: Builds B trees from bootstrap samples
- Trees can be grown in parallel
- Each tree tries to model the full outcome y
- Trees are grown large, left unpruned
- Tuning parameters:
 - B - number of trees (large B is good, helps reduce variance)
 - m - number of predictors considered at each split (small m is good, helps de-correlate the trees)

Boosted trees

- Uses B trees, each fit to the entire training data
- Trees are grown sequentially
- Instead of modeling full outcome y , at each step a tree is built to model r , the latest residual
- Trees are grown very small, “stumps” ($d = 1$) are popular
- Tuning parameters:
 - B - number of trees (if B is too large, big risk of overfitting)
 - d - Tree depth (“interaction depth”)
 - λ - Shrinkage parameter (“learning rate”)


- Both are examples of *ensemble methods*, which are models that form estimates by combining the results of a bunch of other models

Unsupervised learning methods

- K -means clustering
- Hierarchical clustering
 - Single linkage
 - Complete linkage
 - Average linkage
- Gaussian Mixture Models (“EM” clustering)
- Association rule mining
- Principal components analysis

Gaussian Mixture Modeling vs. K -means

Different cluster analysis results on "mouse" data set:



- GMM's do better on this example because they essentially allow for a data-adaptive notion of *distance* when assigning points to centroids
- i.e., In the original data, we have 2 clumps with small variance, and one clump with large variance
- K -means can't capture this added information
- GMM's say: An observation belongs to C_k if its *variance-adjusted* (i.e., Σ_k -adjusted) distance to μ_k is small

K-means

- Assumes spherical clusters
- Minimizes within-cluster sum of squared distances
- **Hard** assignment of points to clusters

Gaussian Mixture Models

- Explicitly allows for non-spherical clusters (via Σ_k)
- Minimizes a slightly different distance, which accounts for priors $\hat{\pi}_k$ and covariances $\hat{\Sigma}_k$
- **Soft** (probabilistic) assignment of points to clusters

- Both require K to be specified in advance, results can change a lot across different choices of K
- Both are fit with very similar looking iterative algorithms
- Algorithms converge to something, but not necessarily the **global optimum**
- Results depend on random initialization
 - Try different random initializations, and keep the best result

Hierarchical clustering

- We discussed **agglomerative** (“bottom-up”) hierarchical clustering
- Hierarchical clustering is performed using a **dissimilarity metric** $d(x_i, x_{i'})$ telling us how dissimilar two **points** are, and a linkage $d(G, H)$ telling us how dissimilar two **clusters** are


<i>Linkage</i>	<i>Description</i>
Complete	Maximal inter-cluster dissimilarity. Compute all pairwise dissimilarities between the observations in cluster A and the observations in cluster B, and record the <i>largest</i> of these dissimilarities.
Single	Minimal inter-cluster dissimilarity. Compute all pairwise dissimilarities between the observations in cluster A and the observations in cluster B, and record the <i>smallest</i> of these dissimilarities.
Average	Mean inter-cluster dissimilarity. Compute all pairwise dissimilarities between the observations in cluster A and the observations in cluster B, and record the <i>average</i> of these dissimilarities.

Single linkage

In **single linkage** (i.e., nearest-neighbor linkage), the dissimilarity between G, H is the smallest dissimilarity between two points in different groups:



$$d_{\text{single}}(G, H) = \min_{i \in G, j \in H} d(x_i, x_j)$$

Example (dissimilarities d_{ij} are distances, groups are marked by colors): single linkage score $d_{\text{single}}(G, H)$ is the distance of the **closest pair**



Single linkage example

Here $n = 60$, $x_i \in \mathbb{R}^2$, $d_{ij} = \|x_i - x_j\|_2$. Cutting the tree at $h = 0.9$ gives the clustering assignments marked by colors




Cut interpretation: for each point x_i , there is another point x_j in its cluster such that $d(x_i, x_j) \leq 0.9$

Complete linkage

In **complete linkage** (i.e., furthest-neighbor linkage), dissimilarity between G, H is the largest dissimilarity between two points in different groups:



$$d_{\text{complete}}(G, H) = \max_{i \in G, j \in H} d(x_i, x_j)$$

Example (dissimilarities d_{ij} are distances, groups are marked by colors): complete linkage score $d_{\text{complete}}(G, H)$ is the distance of the **furthest pair**



Complete linkage example

Same data as before. Cutting the tree at $h = 5$ gives the clustering assignments marked by colors



Cut interpretation: for each point x_i , every other point x_j in its cluster satisfies $d(x_i, x_j) \leq 5$


Average linkage

In **average linkage**, the dissimilarity between G, H is the average dissimilarity over all points in opposite groups:

$$d_{\text{average}}(G, H) = \frac{1}{|G| \cdot |H|} \sum_{i \in G, j \in H} d(x_i, x_j)$$



Example (dissimilarities d_{ij} are distances, groups are marked by colors): average linkage score $d_{\text{average}}(G, H)$ is the **average distance** across all pairs

(Plot here only shows distances between the alert points and one orange point)



Average linkage example

Same data as before. Cutting the tree at $h = 2.5$ gives clustering assignments marked by the colors




Cut interpretation: there really isn't a good one! ☹️

Shortcomings of Single and Complete linkage

Single and complete linkage have some practical problems:

- **Single linkage** suffers from **chaining**.
 - In order to merge two groups, only need one pair of points to be close, irrespective of all others. Therefore *clusters can be too spread out*, and not compact enough.
- **Complete linkage** avoids chaining, but suffers from **crowding**.
 - Because its score is based on the worst-case dissimilarity between pairs, *a point can be closer to points in other clusters than to points in its own cluster*. Clusters are compact, but not far enough apart.
- **Average linkage** tries to **strike a balance**. It uses average pairwise dissimilarity, so clusters tend to be relatively compact and relatively far apart
 - **But**, average linkage doesn't give us a nice interpretation when we cut the dendrogram at height h , and the results change if we apply even a *monotone increasing transformation* to the dissimilarity: E.g., $d \rightarrow d^2$ or $d \rightarrow \frac{e^d}{1+e^d}$

Dimensionality reduction via PCA: A picture



[source: <https://onlinecourses.science.psu.edu/stat857/>]


- This Figure shows an observation $x_i = (x_{i1}, x_{i2})$ along with z_{i1} , its projection onto the first principal component direction, and z_{i2} , its projection onto the second principal component direction

Biplot of US Arrests data

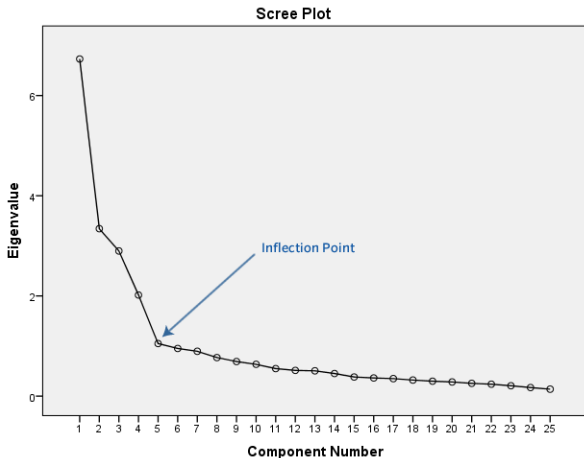
```
arrests.pca$loadings = -arrests.pca$loadings # Flip loadings (phi's)
```

```
arrests.pca$scores = -arrests.pca$scores # Flip scores (z's)
```

```
biplot(arrests.pca, scale = 0) # Construct biplot
```



Selecting number of PCA components: Finding elbows in scree plots



[source: <https://gugginotes.wordpress.com/>]

- *Eigenvalue* *y*-axis label should be interpreted as PVE
- Rule-of-thumb: Stop at the *elbow* in the Scree plot. ($k = 5$ here)

Other things you should know

- Using principal components as variables in regression is called Principal Components Regress (PCR)
- PCR can fail if the top principal components aren't associated with Y
- Correspondence Analysis extends PCA to handle categorical variables
- Multidimensional Scaling (MDS) produces low-dimensional projections of data based on dissimilarity matrices

Cross-validation: Right and Wrong

- Suppose we're in a classification setting with $Y \in \{0, 1\}$
- Getting data points is really *expensive*, so we wind up with only $n = 120$ observations
- However, we have *a lot* of features collected on each observation:
 $p = 5000$
- Fitting a model with $n = 120$ and $p = 5000$ is hopeless, so we take the following approach:
 - ❶ For each of the $p = 5000$ features, calculate the feature's **correlation with y** . Identify the $q = 100$ predictors that have the **highest correlation** with the outcome y .
 - ❷ Apply a classifier using just these 100 predictors
- **Question:** How do we estimate the **test error** of our classifier?
- Can we apply Cross-validation to just Step 2?

Cross-validation: The wrong way

- ❶ For each of the $p = 5000$ features, calculate the feature's correlation with y . Identify the $q = 100$ predictors that have the highest correlation with the outcome y .
- ❷ Apply a classifier using just these 100 predictors
 - It is **very wrong** to estimate test error by applying Cross-validation just on Step 2
 - Doing so would ignore the fact that we used the y values of the training data in Step 1
 - To get valid error estimates, each step of Cross-validation **must** go through **every part** of the model-fitting process that in any way **makes use of the outcome variable y**
 - If we just Cross-validate Step 2, we can **easily** wind up with a CV error estimate of 0% when the true test error is 50%.

WRONG approach

Step 1

Identify the 100 predictors with the highest value of $\text{corr}(X_j, y)$

Cross-validation

Split data into K folds

For $k = 1, 2, \dots, K$:

- Fit classifier using all the 100 selected predictors on all observations except those in fold k
- Use classifier to predict labels for data in fold k
- Calculate test error on fold k

Return: average test error across all of the folds

RIGHT approach

Cross-validation

Split data into K folds

For $k = 1, 2, \dots, K$:

- Using observations not in fold k , identify the 100 predictors with the highest value of $\text{corr}(X_j, y)$
- Fit classifier using these 100 selected predictors on all observations except those in fold k
- Use classifier to predict labels for data in fold k
- Calculate test error on fold k

Return: average test error across all of the folds

Classification with imbalanced data

- Suppose that we're trying to identify *fraudulent transactions*
- We build a random forest classifier, and find that it gets 97% accuracy
- Looking back at the data, we see that just 3% of our observations were from fraudulent transactions
- Our random forest got 97% accuracy just by classifying $\hat{y}_i = \text{not fraudulent}$ for every observation
- This is an example of imbalanced data
- Most classification problems you'll encounter in the real world are imbalanced in this way

Strategies for dealing with imbalanced data

- Stop using *classification accuracy* as your performance metric
 - Instead, assess how well your classifier is doing using more **task-relevant metrics**: Precision, Recall (Sensitivity), Specificity, PPV, **profit**
- Take steps to artificially **improve balance**:
 - Randomly remove instances of your over-represented class (**under-sample**)
 - Randomly add copies of your under-represented class (**over-sample**, sampling *with* replacement)
 - Cluster your over-represented observations into a large number of clusters, and select just one representative observation from each class.
 - Create synthetic samples from your under-represented class using a method like SMOTE (Synthetic Minority Over-sampling Technique)
- If using trees or forests: Try growing very **deep** trees and applying a cost-sensitive pruning technique

How do we approach a particular data analytic task?

In the next few slides I'm going to discuss general “rules of thumb” for approaching a data analytic tasks.

!!Disclaimer!!: There are many reasonable ways to approach problems. The suggestions made in the subsequent slides are at best a helpful over-simplification.

How do we approach a particular data analytic task?

Here's a sequence of questions that you should ask yourself when trying to decide what method(s) to use.

- Is it a supervised or unsupervised learning problem?
- **Unsupervised:**
 - Are trying to group together the observations (cluster), group together the features (cluster) or construct a low-dimensional representation of our data (dimensionality reduction)?
 - If it's a clustering problem:
 - Is the data set very large?
 - > K -means computation scales as $O(n)$, hierarchical clustering is generally $O(n^3)$, so hierclust may not work on large data sets
 - How many clusters do we expect? What shape do we expect/want them to take?
 - > Hierarchical clustering allows us to consider different choices of K in a manner that produces easily comparable clusterings from one K to the next
 - > K -means and GMM's can give very different clusterings for different choices of K

How do we approach a particular data analytic task?

- **Unsupervised:**

- Do we want to **scale** the points in any way?
- If it's a **hierarchical clustering problem**:
 - What kinds of similarity do we want our clustering to capture?
 - > Use a dissimilarity metric that reflects this.
 - Do we need a guarantee on the “connectedness” of our clusters when we cut the dendrogram at level h ? If yes, choose linkage accordingly.
 - > **Single linkage**: Ensures that for every cluster C_k and every $x_i \in C_k$, there's at least one other point $x_{i'} \in C_k$ with $d(x_i, x_{i'}) \leq h$
 - > **Complete linkage**: Ensures that for every cluster C_k , and every pair of points $x_i, x_{i'} \in C_k$, we have $d(x_i, x_{i'}) \leq h$
 - Other good/bad aspects of the 3 linkages we discussed are covered in earlier slides
- If it could be a K -means or GMM problem:
 - Is n sufficiently large and p sufficiently small to be able to estimate the covariance matrices Σ_k for GMM's?
 - Might we have ellipsoidal clusters or spherical clusters where some are more diffuse than others? (GMM's can work better here)
 - Do we want a *hard clustering* (each point gets assigned a cluster), or a *soft clustering* (get an estimated probability that point i is in cluster C_k)

How do we approach a particular data analytic task?

- **Supervised:**

- Should you use Classification or Regression? Might both work?
- How many observations are there?
- How many features are there?
- Is the sample size much larger than the number of features?
- Do I need a very interpretable model?
- Are there many highly correlated predictors?
- Are most of the features likely to be relevant?
- For Classification in particular:
 - What's more costly: False Positives or False Negatives?
 - What is the prevalence of the interesting class? Is the data imbalanced?
 - Do we need good probability estimates, or do we just care about getting the right class labels?

Problem: Small n , large p

- If we have a lot of features and not very many observations, we can try:
 - Variable selection
 - Regularized regression
 - Cluster features or run PCA to reduce dimension
 - Feature screening
 - E.g., Calculate correlations $\text{cor}(X_j, y)$, and keep only the features that have high marginal correlation with y
 - Naive Bayes
- We should **avoid** methods that tend to have high variability (HV), or which require a lot of parameters to be estimated (LP). E.g., avoid:
 - (HV) Large decision trees (Only use pruned trees!)
 - (HV) Unregularized regression
 - (HV) k -Nearest Neighbours
 - (LP) QDA
- Bagging, Random forests and Boosting may still give good results. Proceed with caution, and carefully apply Cross-validation to avoid over-fitting.

Problem: Complex relationship between y and inputs

- Assuming you have the sample size to reliably fit flexible models, try:
 - Decision trees
 - Bagging, Random forests, Boosting
 - LDA, QDA (for classification)
 - Additive models
 - k -Nearest Neighbours
 - k -NN can work well if p is quite small, or if you carefully construct a distance metric that measures distance in just the *relevant directions*
- If you believe there may be important **interaction effects** among the inputs, **avoid**:
 - Additive models
 - Naive Bayes
 - Logistic regression (without interactions)

Acknowledgements

All of the lectures notes for this class feature content borrowed with or without modification from the following sources:

- 36-462/36-662 Lecture notes (Prof. Tibshirani, Prof. G'Sell, Prof. Shalizi)
- 95-791 Lecture notes (Prof. Dubrawski)
- *An Introduction to Statistical Learning, with applications in R* (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani
- *Applied Predictive Modeling*, (Springer, 2013), Max Kuhn and Kjell Johnson